

2.3 - Classify images

- [Overview](#)
- [Introduction](#)
- [Provision Azure resources for Azure AI Custom Vision](#)
- [Understand image classification](#)
- [Train an image classifier](#)
- [Exercise - Classify images with Azure AI Custom Vision](#)
 - [Clone the repository for this course](#)
- [Create Custom Vision resources](#)
 - [Create a Custom Vision project](#)
 - [Test the model](#)
 - [View the project settings](#)
 - [Use the _training_ API](#)
 - [Publish the image classification model](#)
 - [Use the image classifier from a client application](#)
- [Knowledge Check](#)
- [Summary](#)

Overview

Image classification is used to determine the main subject of an image. You can use the **Azure AI Custom Vision** services to train a model that classifies images based on your own categorizations.

Introduction

Image classification is a common computer vision problem that requires software to analyze an image in order to categorize (or *classify*) it. In this module, you will learn how the **Azure AI Custom Vision** service enables you to build your own computer vision models for image classification.


After completing this module, you'll be able to:

- Provision Azure resources for Azure AI Custom Vision.
 - Describe image classification.
 - Train an image classifier.
-

Provision Azure resources for Azure AI Custom Vision

The **Azure AI Custom Vision** service enables you to build your own computer vision models for *image classification* or *object detection*.

Creating an Azure AI Custom Vision solution involves two tasks:



1. Use existing (labeled) images to train an Azure AI Custom Vision model.
2. Create a client application that submits new images to your model to generate predictions.

To use the Azure AI Custom Vision service, you must provision two kinds of Azure resource:

- A *training* resource used to train your models. This can be:
 - An **Azure AI services multi-service** resource.
 - An **Azure AI Custom Vision (Training)** resource.
- A *prediction* resource, used by client applications to get predictions from your model. This can be:
 - An **Azure AI services multi-service** resource.
 - An **Azure AI Custom Vision (Prediction)** resource.

You can use a **Azure AI services multi-service** resource for both training and prediction, and you can mix-and-match resource types (for example, using an **Azure AI Custom Vision (Training)** resource to train a model that you then publish using an **Azure AI services multi-service** resource). If using a multi-service resource, the key and endpoint for both training and prediction will be the same.

Understand image classification

Image classification is a computer vision technique in which a model is trained to predict a class label for an image based on its contents. Usually, the class label relates to the main *subject* of the image.

For example, the following images have been classified based on the type of fruit they contain.



Apple



Banana



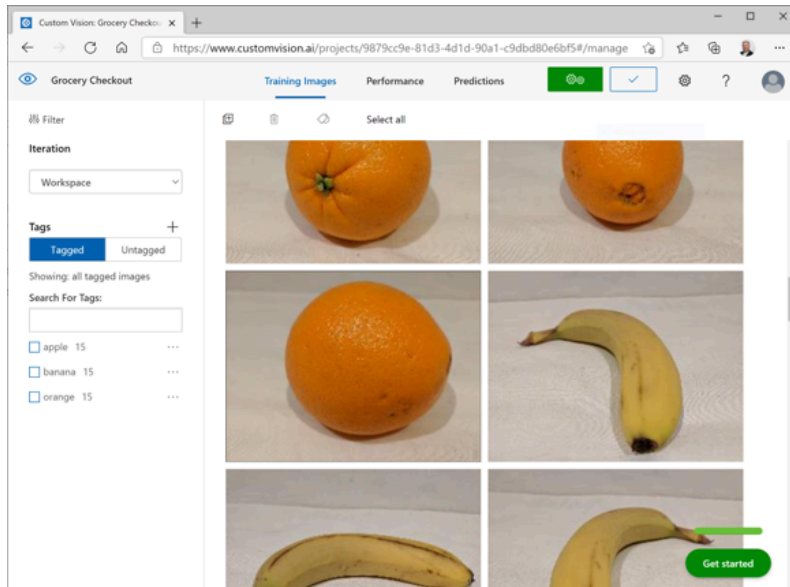
Orange

Models can be trained for multiclass classification (in other words, there are multiple classes, but each image can belong to only one class) or multilabel classification (in other words, an image might be associated with multiple labels).

Train an image classifier

To train an image classification model with the Azure AI Custom Vision service, you can use the Azure AI Custom Vision portal, the Azure AI Custom Vision REST API or SDK, or a combination of both approaches.

In most cases, you'll typically use the Azure AI Custom Vision portal to train your model.



The portal provides a graphical interface that you can use to:

1. Create an image classification project for your model and associate it with a training resource.
2. Upload images, assigning class label tags to them.
3. Review and edit tagged images.
4. Train and evaluate a classification model.
5. Test a trained model.
6. Publish a trained model to a prediction resource.

The REST API and SDKs enable you to perform the same tasks by writing code, which is useful if you need to automate model training and publishing as part of a DevOps process.

Exercise - Classify images with Azure AI Custom Vision

The **Azure AI Custom Vision** service enables you to create computer vision models that are trained on your own images. You can use it to train *image classification* and *object detection* models; which you can then publish and consume from applications.

In this exercise, you will use the Custom Vision service to train an image classification model that can identify three classes of fruit (apple, banana, and orange).

Clone the repository for this course

If you have not already cloned **mslearn-ai-vision** code repository to the environment where you're working on this lab, follow these steps to do so. Otherwise, open the cloned folder in Visual Studio Code.

1. Start Visual Studio Code.
2. Open the palette (SHIFT+CTRL+P) and run a **Git: Clone** command to clone the <https://github.com/MicrosoftLearning/mslearn-ai-vision> repository to a local folder (it doesn't matter which folder).
3. When the repository has been cloned, open the folder in Visual Studio Code.
4. Wait while additional files are installed to support the C# code projects in the repo.

Note: If you are prompted to add required assets to build and debug, select **Not Now**. If you are prompted with the Message *Detected an Azure Function Project in folder*, you can safely close that message.

Create Custom Vision resources

Before you can train a model, you will need Azure resources for *training* and *prediction*. You can create **Custom Vision** resources for each of these tasks, or you can create a single **Azure AI Services** resource and use it for either (or both).

In this exercise, you'll create **Custom Vision** resources for training and prediction so that you can manage access and costs for these workloads separately.

1. In a new browser tab, open the Azure portal at <https://portal.azure.com>, and sign in using the Microsoft account associated with your Azure subscription.
2. Select the **+ Create a resource** button, search for *custom vision*, and create a **Custom Vision** resource with the following settings:
 - **Create options:** Both
 - **Subscription:** *Your Azure subscription*
 - **Resource group:** *Choose or create a resource group (if you are using a restricted subscription, you may not have permission to create a new resource group - use the one provided)*
 - **Region:** *Choose any available region*
 - **Name:** *Enter a unique name*
 - **Training pricing tier:** F0
 - **Prediction pricing tier:** F0

Note: If you already have an F0 custom vision service in your subscription, select **S0** for this one.

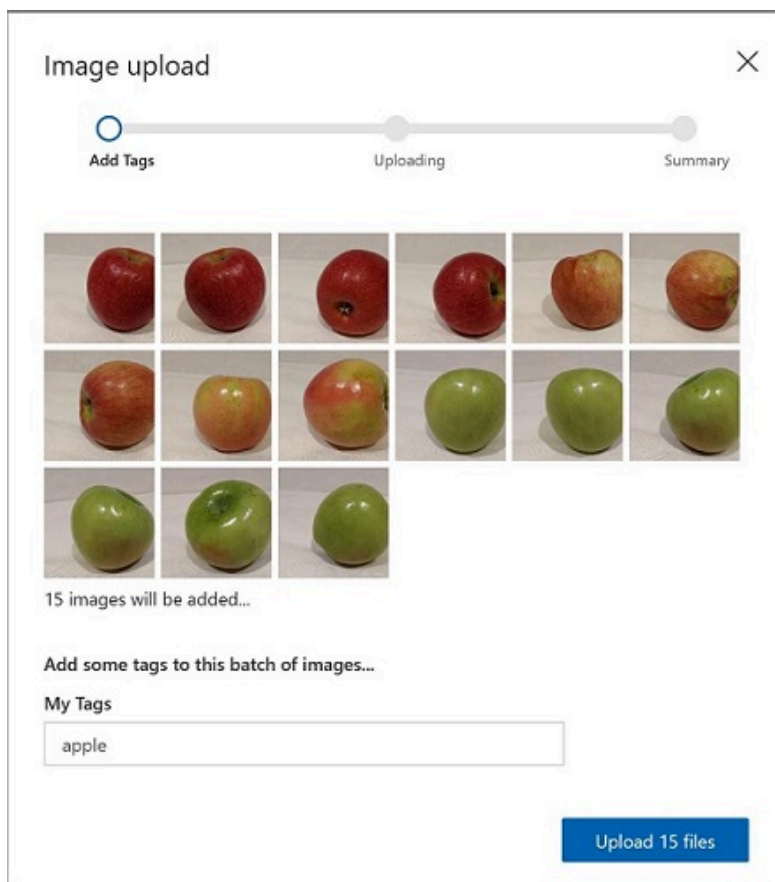
3. Wait for the resources to be created, and then view the deployment details and note that two Custom Vision resources are provisioned; one for training, and another for prediction (evident by the - **Prediction** suffix). You can view these by navigating to the resource group where you created them.

Important: Each resource has its own *endpoint* and *keys*, which are used to manage access from your code. To train an image classification model, your code must use the *training* resource (with its endpoint and key); and to use the trained model to predict image classes, your code must use the *prediction* resource (with its endpoint and key).

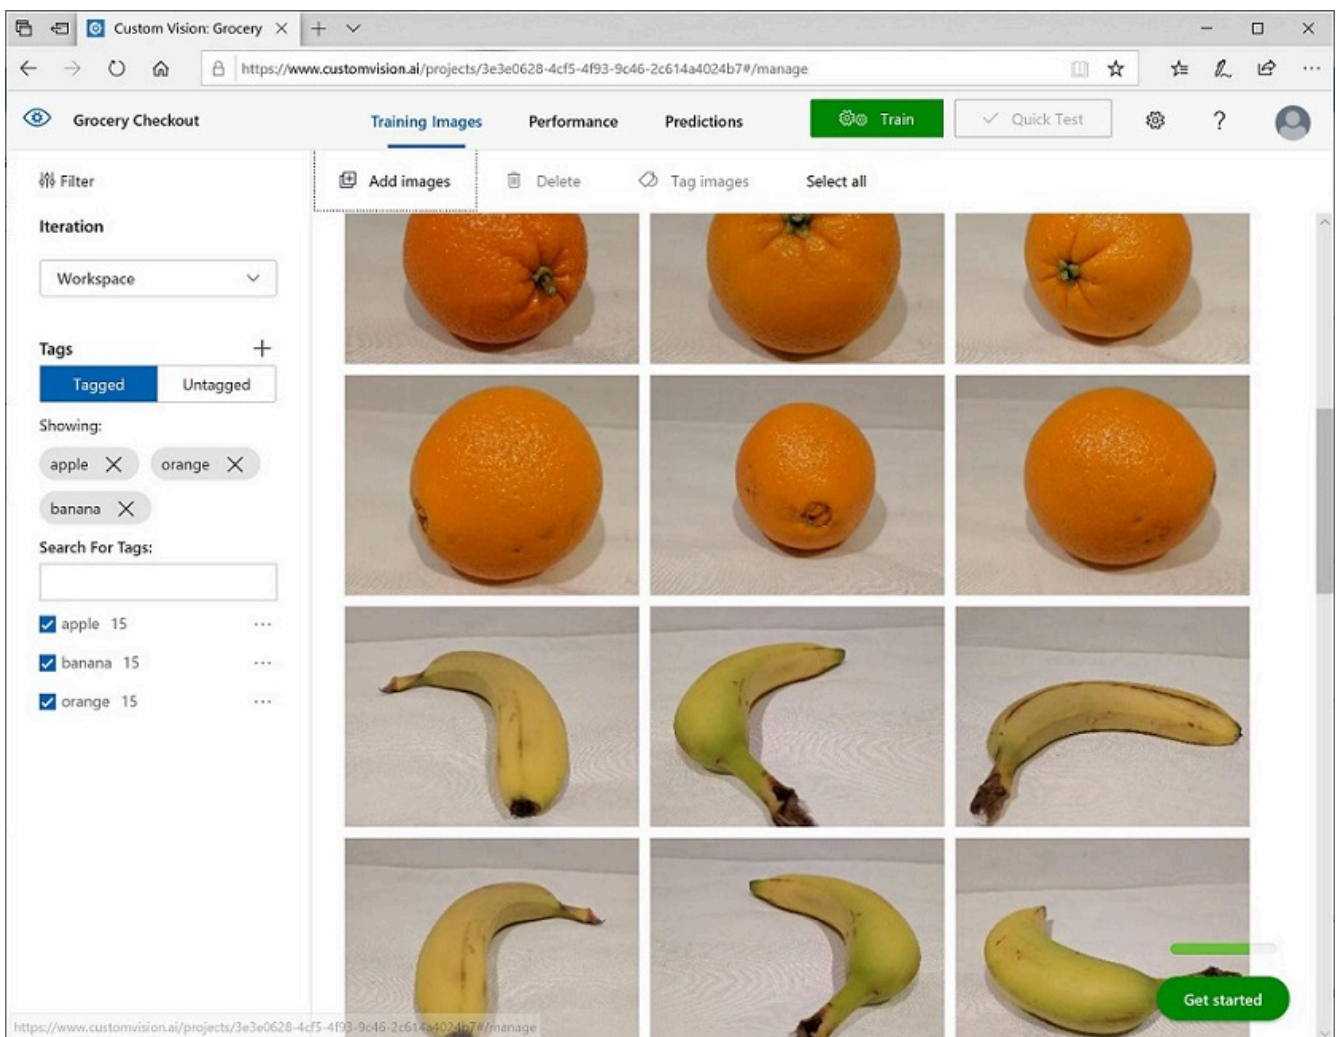
Create a Custom Vision project

To train an image classification model, you need to create a Custom Vision project based on your training resource. To do this, you'll use the Custom Vision portal.

1. In Visual Studio Code, view the training images in the **07-custom-vision-image-classification/training-images** folder where you cloned the repository. This folder contains subfolders of apple, banana, and orange images.
2. In a new browser tab, open the Custom Vision portal at <https://customvision.ai>. If prompted, sign in using the Microsoft account associated with your Azure subscription and agree to the terms of service.
3. In the Custom Vision portal, create a new project with the following settings:
 - **Name:** Classify Fruit
 - **Description:** Image classification for fruit
 - **Resource:** *The Custom Vision resource you created previously*
 - **Project Types:** Classification
 - **Classification Types:** Multiclass (single tag per image)
 - **Domains:** Food
4. In the new project, click **[+] Add images**, and select all of the files in the **training-images/apple** folder you viewed previously. Then upload the image files, specifying the tag *apple*, like this:



1. Repeat the previous step to upload the images in the **banana** folder with the tag *banana*, and the images in the **orange** folder with the tag *orange*.
2. Explore the images you have uploaded in the Custom Vision project - there should be 15 images of each class, like this:



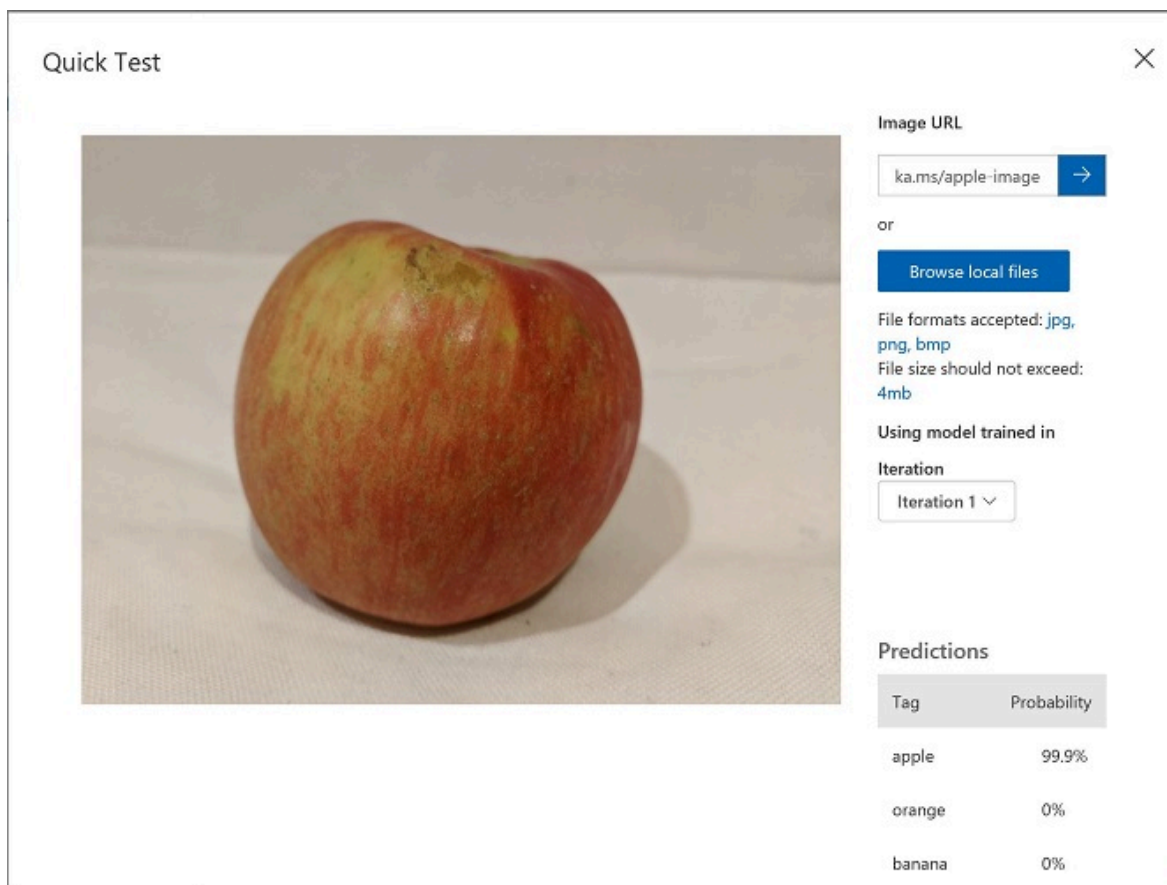
1. In the Custom Vision project, above the images, click **Train** to train a classification model using the tagged images. Select the **Quick Training** option, and then wait for the training iteration to complete (this may take a minute or so).
2. When the model iteration has been trained, review the *Precision*, *Recall*, and *AP* performance metrics - these measure the prediction accuracy of the classification model, and should all be high.

Note: The performance metrics are based on a **probability threshold of 50% for each prediction** (in other words, if the model calculates a 50% or higher probability that an image is of a particular class, then that class is predicted). You can adjust this at the top-left of the page.

Test the model

Now that you've trained the model, you can test it.

1. Above the performance metrics, click **Quick Test**.
2. In the **Image URL** box, type `https://aka.ms/apple-image` and click →
3. View the predictions returned by your model - the probability score for *apple* should be the highest, like this:



1. Close the **Quick Test** window.

View the project settings

The project you have created has been assigned a unique identifier, which you will need to specify in any code that interacts with it.

1. Click the *settings* (⚙️) icon at the top right of the **Performance** page to view the project settings.
2. Under **General** (on the left), note the **Project Id** that uniquely identifies this project.
3. On the right, under **Resources** note that the key and endpoint are shown. These are the details for the *training* resource (you can also obtain this information by viewing the resource in the Azure portal).

Use the *training* API

The Custom Vision portal provides a convenient user interface that you can use to upload and tag images, and train models. However, in some scenarios you may want to automate model training by using the Custom Vision training API.

Note: In this exercise, you can choose to use the API from either the **C#** or **Python** SDK. In the steps below, perform the actions appropriate for your preferred language.

1. In Visual Studio Code, in the **Explorer** pane, browse to the **07-custom-vision-image-classification** folder and expand the **C-Sharp** or **Python** folder depending on your language preference.
2. Right-click the **train-classifier** folder and open an integrated terminal. Then install the Custom Vision Training package by running the appropriate command for your language preference:

C#

```
dotnet add package Microsoft.Azure.CognitiveServices.Vision.CustomVision.Training --version 2.0.0
```

Python

```
pip install azure-cognitiveservices-vision-customvision==3.1.0
```

1. View the contents of the **train-classifier** folder, and note that it contains a file for configuration settings:

- **C#:** appsettings.json
- **Python:** .env

Open the configuration file and update the configuration values it contains to reflect the endpoint and key for your Custom Vision *training* resource, and the project ID for the classification project you created previously. Save your changes.

2. Note that the **train-classifier** folder contains a code file for the client application:

- **C#:** Program.cs
- **Python:** [train-classifier.py](#)

Open the code file and review the code it contains, noting the following details:

- Namespaces from the package you installed are imported
- The **Main** function retrieves the configuration settings, and uses the key and endpoint to create an authenticated **CustomVisionTrainingClient**, which is then used with the project ID to create a **Project** reference to your project.
- The **Upload_Images** function retrieves the tags that are defined in the Custom Vision project and then uploads image files from correspondingly named folders to the project, assigning the appropriate tag ID.
- The **Train_Model** function creates a new training iteration for the project and waits for training to complete.

3. Return the integrated terminal for the **train-classifier** folder, and enter the following command to run the program:

C#

```
dotnet run
```

Python

```
python train-classifier.py
```

1. Wait for the program to end. Then return to your browser and view the **Training Images** page for your project in the Custom Vision portal (refreshing the browser if necessary).
2. Verify that some new tagged images have been added to the project. Then view the **Performance** page and verify that a new iteration has been created.

Publish the image classification model

Now you're ready to publish your trained model so that it can be used from a client application.

1. In the Custom Vision portal, on the **Performance** page, click ✓ **Publish** to publish the trained model with the following settings:
 - **Model name:** fruit-classifier
 - **Prediction Resource:** *The **prediction** resource you created previously which ends with “-Prediction” (not the training resource).*
2. At the top left of the **Project Settings** page, click the *Projects Gallery* (👁) icon to return to the Custom Vision portal home page, where your project is now listed.
3. On the Custom Vision portal home page, at the top right, click the *settings* (⚙) icon to view the settings for your Custom Vision service.
4. Then, under **Resources**, find your *prediction* resource which ends with “-Prediction” (not the training resource) to determine its **Key** and **Endpoint** values (you can also obtain this information by viewing the resource in the Azure portal).

Use the image classifier from a client application

Now that you’ve published the image classification model, you can use it from a client application. Once again, you can choose to use **C#** or **Python**.

1. In Visual Studio Code, in the **07-custom-vision-image-classification** folder, in the subfolder for your preferred language (**C-Sharp** or **Python**), right-click the **test-classifier** folder and open an integrated terminal. Then enter the following SDK-specific command to install the Custom Vision Prediction package:

C#

```
dotnet add package Microsoft.Azure.CognitiveServices.Vision.CustomVision.Prediction --version 2.0.0
```

Python

```
pip install azure-cognitiveservices-vision-customvision==3.1.0
```

Note: The Python SDK package includes both training and prediction packages, and may already be installed.

1. Expand the **test-classifier** folder to view the files it contains, which are used to implement a test client application for your image classification model.
2. Open the configuration file for your client application (*appsettings.json* for C# or *.env* for Python) and update the configuration values it contains to reflect the endpoint and key for your Custom Vision *prediction* resource, the project ID for the classification project, and the name of your published model (which should be *fruit-classifier*). Save your changes.
3. Open the code file for your client application (*Program.cs* for C#, *test-classifier.py* for Python) and review the code it contains, noting the following details:
 - Namespaces from the package you installed are imported
 - The **Main** function retrieves the configuration settings, and uses the key and endpoint to create an authenticated **CustomVisionPredictionClient**.
 - The prediction client object is used to predict a class for each image in the **test-images** folder, specifying the project ID and model name for each request. Each prediction includes a probability for each possible class, and only predicted tags with a probability greater than 50% are displayed.

- View the file here: [test-classifier.py](#)

4. Return the integrated terminal for the **test-classifier** folder, and enter the following SDK-specific command to run the program:

Note: You can publish a trained Azure AI Custom Vision model to either **an Azure AI Custom Vision (Prediction) resource** or **an Azure AI Services multi-service resource**.

C#

```
dotnet run
```

Python

```
python test-classifier.py
```

1. View the label (tag) and probability scores for each prediction. You can view the images in the **test-images** folder to verify that the model has classified them correctly.

Knowledge Check

1. You want to train a model that can categorize an image as "cat" or "dog" based on its subject. What kind of Azure AI Custom Vision project should you create? *

☒ Image classification (multiclass)

✓ That's correct. To train a model that classifies an image using a single tag, use an Image classification (multiclass) project.

☐ Image classification (multilabel)

☐ Object detection

2. Which of the following kinds of Azure resource can you use to host a trained Azure AI Custom Vision model? *

☐ Azure AI Custom Vision (Training)

☐ Azure AI Vision

☒ Azure AI Services

✓ That's correct. You can publish a trained Azure AI Custom Vision model to either an Azure AI Custom Vision (Prediction) resource or an Azure AI Services multi-service resource.

Summary

In this module, you learned how the **Azure AI Custom Vision** service enables you to build your own computer vision models for image classification.

Now that you've completed this module, you can:

- Provision Azure resources for Azure AI Custom Vision.
- Describe image classification.
- Train an image classifier.

To find out more about the Azure AI Custom Vision service, see the [Azure AI Custom Vision documentation](#).

👉 Compiled by [Kenneth Leung](#) (2025)